

Estratégias Algorítmicas 2020/21

Week 11 – Graph Algorithms (SCC and MaxFlow)




UNIVERSIDADE DE COIMBRA

1. Strongly Connected Components
2. Max-flow Algorithms
 - 2.1 Ford-Fulkerson Algorithm
 - 2.2 Edmond-Karp Algorithm
 - 2.3 Applications

Introduction


Given a directed graph $G = (V, A)$, a subgraph G' is a **strongly connected component** if there exists a path between a vertex and every other vertex in G' and this subgraph has maximal size.




3 strongly connected components

Strongly Connected Components

A vertex v is the root of a connected component if $low[v] = dfs[v]$



A graph



DFS tree

The vertices in each strongly connected component under the root are stored in a stack. It is possible to solve it with Tarjan Algorithm in $O(|V| + |A|)$.

Strongly Connected Components

Function *Tarjan*(v)

$low[v] = dfs[v] = t$

$t = t + 1$

push(S, v)

for each arc $(v, w) \in A$ **do**

if $dfs[w]$ has no value **then**

Tarjan(w)

$low[v] = \min(low[v], low[w])$

else if $w \in S$ **then**

$low[v] = \min(low[v], dfs[w])$

if $low[v] = dfs[v]$ **then**

$C = \emptyset$

repeat

$w = pop(S)$

push(C, w)

until $w = v$

push(Scc, C)

Scc collects all stacks of strongly connected components. This algorithm must be called for every unvisited vertex in the graph.

Introduction


Given a directed graph $G = (V, A)$, where each arc (u, v) has a capacity $c(u, v) > 0$ (a *network*), and two vertices *source* s and *sink* t , the **maximum flow problem** consists of maximizing the total amount of flow from s to t subject to two constraints:

1. Flow on arc (u, v) does not exceed $c(u, v)$
2. For every vertex $v \neq s, t$, incoming flow is equal to outgoing flow.


This problem arises in many real-life situations:

- Routing as many packets as possible on a network, given a bandwidth capacity
- Sending as many trucks as possible, where roads have limits on the number of trucks per unit time

Introduction



A network



A network with maximum flow of 6

How to compute the maximum flow?


Ford-Fulkerson Algorithm

A residual network $G_R = (V, A_R)$ has the same vertices as the original network G and two different types of arcs:


1. For each arc (u, v) with capacity $c(u, v)$ and flow $f(u, v) < c(u, v)$, there exists an arc (u, v) in the residual network with capacity $c(u, v) - f(u, v)$
2. For each arc (u, v) with capacity $c(u, v)$ and flow $f(u, v) > 0$, there exists an arc (v, u) in the residual network with capacity $f(u, v)$

An augmenting path is an path from source to sink in the residual network G^* that should send as much flow as possible.

Ford-Fulkerson Algorithm



Network flow



Residual network

Ford-Fulkerson Algorithm

Function $FF(G_R)$

$mflow = 0$

while there is a s - t path P in G_R **do**

$f_p = \min\{c(u, v) \mid (u, v) \in P\}$

for each arc (u, v) in P **do**

$c(u, v) = c(u, v) - f_p$


$c(v, u) = c(v, u) + f_p$

$mflow = mflow + f_p$


return $mflow$

For each arc (u, v) in G , G_R contains an arc (u, v) with capacity $c(u, v)$ and an arc (v, u) with capacity 0.

Ford-Fulkerson Algorithm




Network flow




Residual network

Ford-Fulkerson Algorithm




Network flow




Residual network

Send 1 unit of flow through (s, b, c, t) in the residual network

Ford-Fulkerson Algorithm




Network flow




Residual network

Send 2 units of flow through (s, a, e, t) in the residual network

Ford-Fulkerson Algorithm




Network flow




Residual network

Send 1 unit of flow through (s, c, a, e, t) in the residual network

Ford-Fulkerson Algorithm




Network flow




Residual network

Send 1 unit of flow through (s, b, t) in the residual network

Ford-Fulkerson Algorithm



Network flow




Residual network

Send 1 unit of flow through (s, c, b, t) in the residual network, which cancels the flow from c to b in the network

Edmond-Karp Algorithm

Ford-Fulkerson works under any choice of a s - t path. However, an arbitrary choice may give a bad worst case time complexity.

Consider that the s - t path passes always through vertices u and v :




The number of iterations until it stops is twice the value of the maximum flow.

But if the s - t path with the least number of arcs is chosen (with BFS), the time complexity reduces to $O(|V| \cdot |E|^2)$ for any network ([Edmond-Karp Algorithm](#)).


Variants

Many sources and/or many sinks: Add a source vertex that connects to all sources and/or add a sink vertex to which all sinks are connected. The new arcs must have an "infinite" capacity.





Variants

A vertex a with capacity constraint x : Split the vertex into two, a' and a'' , and connect them by an arc with capacity x .



Minimum Cut Problem


Given a directed graph $G = (V, A)$, where each arc (u, v) has a weight $w(u, v)$, and two vertices s and t , the **minimum cut problem** consists of finding the set of arcs to be removed (a *cut*) with the least total weight such that s and t become unreachable.



Solve a maximum flow problem where the arc weight is the arc capacity. The least total weight is equal to the maximum flow (*max-flow min-cut theorem*).

Arc-Disjoint Paths Problem


Given a directed graph $G = (V, A)$, and two vertices s and t , the **arc-disjoint paths problem** consists of finding the maximum number of s - t paths that do not have arcs in common.




Add capacity one to each arc and solve the corresponding maximum flow problem. The number of arc-disjoint paths is equal to the maximum flow.

Maximum Bipartite Matching

Given a directed graph $G = (U, V, E)$, where U and V are two disjoint sets of vertices, and E is a set of edges, each of which connects an element of U to an element of V , the **maximum bipartite matching problem** consists of finding the maximum number of edges such that no two edges share a vertex (a *matching*).



A bipartite graph




Transformed graph


Transformed graph: add a source that connects to all vertices in U and a sink to which all vertices in V are connected. Arc capacity is equal to 1.

Maximum Bipartite Matching

Given a directed graph $G = (U, V, E)$, where U and V are two disjoint sets of vertices, and E is a set of edges, each of which connects an element of U to an element of V , the **maximum bipartite matching problem** consists of finding the maximum number of edges such that no two edges share a vertex (a *matching*).



A bipartite graph



Transformed graph

The number of edges in a maximal bipartite matching is equal to the maximum flow in the transformed graph